

In [1]:

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5
6
7 pd.set_option('display.max_columns', 80)
```

In [2]:

```
1 df = pd.read_csv('zindi training.csv')
2 df.head(5)
```

Out[2]:

	TransactionId	BatchId	AccountId	SubscriptionId	CustomerId	Curr
0	TransactionId_76871	BatchId_36123	AccountId_3957	SubscriptionId_887	CustomerId_4406	
1	TransactionId_73770	BatchId_15642	AccountId_4841	SubscriptionId_3829	CustomerId_4406	
2	TransactionId_26203	BatchId_53941	AccountId_4229	SubscriptionId_222	CustomerId_4683	
3	TransactionId_380	BatchId_102363	AccountId_648	SubscriptionId_2185	CustomerId_988	
4	TransactionId_28195	BatchId_38780	AccountId_4841	SubscriptionId_3829	CustomerId_988	

In [3]:

```
1 df.shape
```

Out[3]:

```
(95662, 16)
```

In [4]:

1 df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 95662 entries, 0 to 95661
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   TransactionId          95662 non-null  object
1   BatchId                95662 non-null  object
2   AccountId              95662 non-null  object
3   SubscriptionId         95662 non-null  object
4   CustomerId             95662 non-null  object
5   CurrencyCode           95662 non-null  object
6   CountryCode            95662 non-null  int64
7   ProviderId             95662 non-null  object
8   ProductId              95662 non-null  object
9   ProductCategory        95662 non-null  object
10  ChannelId              95662 non-null  object
11  Amount                 95662 non-null  float64
12  Value                  95662 non-null  int64
13  TransactionStartTime    95662 non-null  object
14  PricingStrategy         95662 non-null  int64
15  FraudResult            95662 non-null  int64
dtypes: float64(1), int64(4), object(11)
memory usage: 11.7+ MB
```

In [5]:

```
1 categorical = [var for var in df.columns if df[var].dtype=='O']
2
3 print('There are {} categorical variabes \n'.format(len(categorical)))
4
5 print('They are: ', categorical)
```

There are 11 categorical variabes

They are: ['TransactionId', 'BatchId', 'AccountId', 'SubscriptionId', 'CustomerId', 'CurrencyCode', 'ProviderId', 'ProductId', 'ProductCategory', 'ChannelId', 'TransactionStartTime']

In [6]:

1 df[categorical].head()

Out[6]:

	TransactionId	BatchId	AccountId	SubscriptionId	CustomerId	Cui
0	TransactionId_76871	BatchId_36123	AccountId_3957	SubscriptionId_887	CustomerId_4406	
1	TransactionId_73770	BatchId_15642	AccountId_4841	SubscriptionId_3829	CustomerId_4406	
2	TransactionId_26203	BatchId_53941	AccountId_4229	SubscriptionId_222	CustomerId_4683	
3	TransactionId_380	BatchId_102363	AccountId_648	SubscriptionId_2185	CustomerId_988	
4	TransactionId_28195	BatchId_38780	AccountId_4841	SubscriptionId_3829	CustomerId_988	

In [7]:

```
1 df[categorical].isnull().sum()
```

Out[7]:

```
TransactionId      0
BatchId            0
AccountId          0
SubscriptionId     0
CustomerId         0
CurrencyCode       0
ProviderId         0
ProductId          0
ProductCategory    0
ChannelId          0
TransactionStartTime 0
dtype: int64
```

In [8]:

```
1 for var in categorical:
2     print(var, 'contains', len(df[var].unique()), 'unique values')
```

```
TransactionId contains 95662 unique values
BatchId contains 94809 unique values
AccountId contains 3633 unique values
SubscriptionId contains 3627 unique values
CustomerId contains 3742 unique values
CurrencyCode contains 1 unique values
ProviderId contains 6 unique values
ProductId contains 23 unique values
ProductCategory contains 9 unique values
ChannelId contains 4 unique values
TransactionStartTime contains 94556 unique values
```

In [9]:

```
1 df['TransactionStartTime'] = pd.to_datetime(df['TransactionStartTime'])
```

In [10]:

```
1 df['Year'] = df['TransactionStartTime'].dt.year
2 df['Year'].head()
```

Out[10]:

```
0    2018
1    2018
2    2018
3    2018
4    2018
Name: Year, dtype: int64
```

In [11]:

```
1 df['Month'] = df['TransactionStartTime'].dt.month
2 df['Month'].head()
```

Out[11]:

```
0    11
1    11
2    11
3    11
4    11
```

Name: Month, dtype: int64

In [12]:

```
1 #extract the day
2
3
4 df['day'] = df['TransactionStartTime'].dt.day
5 df['day'].head()
```

Out[12]:

```
0    15
1    15
2    15
3    15
4    15
```

Name: day, dtype: int64

In [13]:

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 95662 entries, 0 to 95661
Data columns (total 19 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   TransactionId                         95662 non-null  object
1   BatchId                              95662 non-null  object
2   AccountId                            95662 non-null  object
3   SubscriptionId                       95662 non-null  object
4   CustomerId                           95662 non-null  object
5   CurrencyCode                         95662 non-null  object
6   CountryCode                          95662 non-null  int64
7   ProviderId                           95662 non-null  object
8   ProductId                            95662 non-null  object
9   ProductCategory                      95662 non-null  object
10  ChannelId                            95662 non-null  object
11  Amount                               95662 non-null  float64
12  Value                               95662 non-null  int64
13  TransactionStartTime                 95662 non-null  datetime64[ns, UTC]
14  PricingStrategy                     95662 non-null  int64
15  FraudResult                         95662 non-null  int64
16  Year                                95662 non-null  int64
17  Month                               95662 non-null  int64
18  day                                 95662 non-null  int64
dtypes: datetime64[ns, UTC](1), float64(1), int64(7), object(10)
memory usage: 13.9+ MB
```

In [14]:

```
1 df.drop('TransactionStartTime',axis = 1, inplace=True)
2 df
```

Out[14]:

	TransactionId	BatchId	AccountId	SubscriptionId	CustomerId
0	TransactionId_76871	BatchId_36123	AccountId_3957	SubscriptionId_887	CustomerId_4406
1	TransactionId_73770	BatchId_15642	AccountId_4841	SubscriptionId_3829	CustomerId_4406
2	TransactionId_26203	BatchId_53941	AccountId_4229	SubscriptionId_222	CustomerId_4683
3	TransactionId_380	BatchId_102363	AccountId_648	SubscriptionId_2185	CustomerId_988
4	TransactionId_28195	BatchId_38780	AccountId_4841	SubscriptionId_3829	CustomerId_988
...
95657	TransactionId_89881	BatchId_96668	AccountId_4841	SubscriptionId_3829	CustomerId_3078
95658	TransactionId_91597	BatchId_3503	AccountId_3439	SubscriptionId_2643	CustomerId_3874
95659	TransactionId_82501	BatchId_118602	AccountId_4841	SubscriptionId_3829	CustomerId_3874
95660	TransactionId_136354	BatchId_70924	AccountId_1346	SubscriptionId_652	CustomerId_1709
95661	TransactionId_35670	BatchId_29317	AccountId_4841	SubscriptionId_3829	CustomerId_1709

95662 rows × 18 columns

In [15]:

```
1 #we'll extraxt the categorical features again since we have modified our data
2 categorical = [var for var in df.columns if df[var].dtype=='O']
3
4 print('There are {} categorical variabes \n'.format(len(categorical)))
5
6 print('They are: ', categorical)
```

There are 10 categorical variabes

They are: ['TransactionId', 'BatchId', 'AccountId', 'SubscriptionId', 'CustomerId', 'CurrencyCode', 'ProviderId', 'ProductId', 'ProductCategory', 'ChannelId']

In [16]:

```
1 df[categorical].isnull().sum()
```

Out[16]:

```
TransactionId      0
BatchId            0
AccountId          0
SubscriptionId     0
CustomerId         0
CurrencyCode       0
ProviderId         0
ProductId          0
ProductCategory    0
ChannelId          0
dtype: int64
```

In [17]:

```
1 len(df.TransactionId.unique())
```

Out[17]:

95662

In [18]:

```
1 #Let's get dummies for this columns
2 providerId_dummy = pd.get_dummies(df.ProviderId, drop_first=True)
3 providerId_dummy
```

Out[18]:

	ProviderId_2	ProviderId_3	ProviderId_4	ProviderId_5	ProviderId_6
0	0	0	0	0	1
1	0	0	1	0	0
2	0	0	0	0	1
3	0	0	0	0	0
4	0	0	1	0	0
...
95657	0	0	1	0	0
95658	0	0	0	0	1
95659	0	0	1	0	0
95660	0	0	0	0	1
95661	0	0	1	0	0

95662 rows × 5 columns

In [19]:

```
1 productId_dummy = pd.get_dummies(df.ProductId, drop_first=True)
2 productId_dummy
```

Out[19]:

	ProductId_10	ProductId_11	ProductId_12	ProductId_13	ProductId_14	ProductId_15	Pr
0	1	0	0	0	0	0	
1	0	0	0	0	0	0	
2	0	0	0	0	0	0	
3	0	0	0	0	0	0	
4	0	0	0	0	0	0	
...	
95657	0	0	0	0	0	0	
95658	1	0	0	0	0	0	
95659	0	0	0	0	0	0	
95660	0	0	0	0	0	0	
95661	0	0	0	0	0	0	

95662 rows × 22 columns

In [20]:

```
1 ProductCategory_dummy = pd.get_dummies(df.ProductCategory, drop_first=True)
2 ProductCategory_dummy
3
```

Out[20]:

	data_bundles	financial_services	movies	other	ticket	transport	tv	utility_bill
0	0	0	0	0	0	0	0	0
1	0	1	0	0	0	0	0	0
2	0	0	0	0	0	0	0	0
3	0	0	0	0	0	0	0	1
4	0	1	0	0	0	0	0	0
...
95657	0	1	0	0	0	0	0	0
95658	0	0	0	0	0	0	0	0
95659	0	1	0	0	0	0	0	0
95660	0	0	0	0	0	0	1	0
95661	0	1	0	0	0	0	0	0

95662 rows × 8 columns

In [21]:

```
1 ChannelId_dummy = pd.get_dummies(df.ChannelId, drop_first=True)
2 ChannelId_dummy
3
```

Out[21]:

	ChannelId_2	ChannelId_3	ChannelId_5
0	0	1	0
1	1	0	0
2	0	1	0
3	0	1	0
4	1	0	0
...
95657	1	0	0
95658	0	1	0
95659	1	0	0
95660	0	1	0
95661	1	0	0

95662 rows × 3 columns

In [22]:

```
1 PricingStrategy_dummy = pd.get_dummies(df.PricingStrategy, drop_first=True)
2 PricingStrategy_dummy
3
```

Out[22]:

	1	2	4
0	0	1	0
1	0	1	0
2	0	1	0
3	0	1	0
4	0	1	0
...
95657	0	1	0
95658	0	1	0
95659	0	1	0
95660	0	1	0
95661	0	1	0

95662 rows × 3 columns

In [23]:

```
1 #we'll extraxt the numerical features
2 numerical = [var for var in df.columns if df[var].dtype!='O']
3
4 print('There are {} numerical variabes \n'.format(len(numerical)))
5
6 print('They are: ', numerical)
```

There are 8 numerical variabes

They are: ['CountryCode', 'Amount', 'Value', 'PricingStrategy', 'FraudResul
t', 'Year', 'Month', 'day']

In [24]:

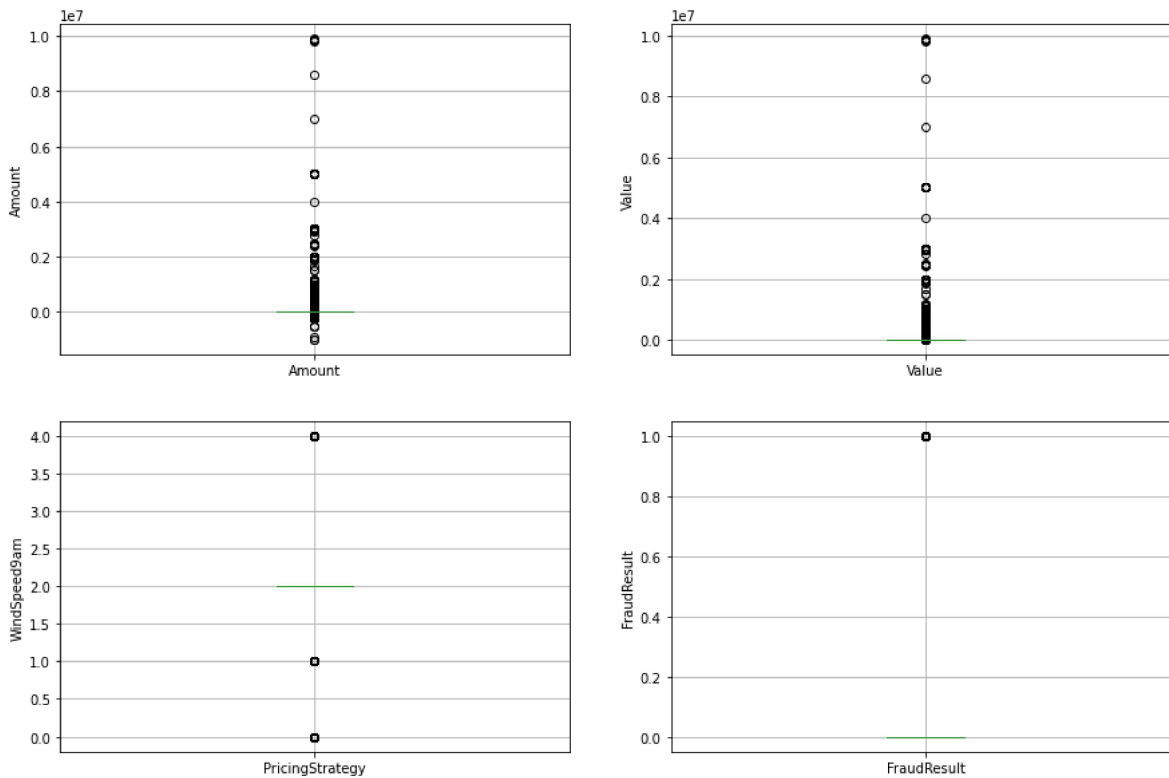
```

1 plt.figure(figsize=(15,10))
2
3 plt.subplot(2,2,1)
4 fig = df.boxplot(column='Amount')
5 fig.set_ylabel('Amount')
6
7 plt.subplot(2,2,2)
8 fig = df.boxplot(column='Value')
9 fig.set_ylabel('Value')
10
11 plt.subplot(2,2,3)
12 fig = df.boxplot(column='PricingStrategy')
13 fig.set_ylabel('WindSpeed9am')
14
15 plt.subplot(2,2,4)
16 fig = df.boxplot(column='FraudResult')
17 fig.set_ylabel('FraudResult')

```

Out[24]:

Text(0, 0.5, 'FraudResult')




In [25]:

```
1 #set figure size
2 plt.figure(figsize=(15,10))
3
4 plt.subplot(2,2,1)
5 fig = df.Amount.hist()
6 fig.set_ylabel('Amount')
7
8 plt.subplot(2,2,2)
9 fig = df.Value.hist()
10 fig.set_ylabel('Value')
11
12
13 plt.subplot(2,2,3)
14 fig = df.PricingStrategy.hist()
15 fig.set_ylabel('PricingStrategy')
16
17 plt.subplot(2,2,4)
18 fig = df.FraudResult.hist()
19 fig.set_ylabel('FraudResult')
```

Out[25]:

Text(0, 0.5, 'FraudResult')



In [26]:

```
1 new_df = df['BatchId'].str.split('_',n=1,expand = True)
2
3 df['BatchId'] = new_df[1].astype(int)
```

In [27]:

```
1 new_df1 = df['CustomerId'].str.split('_',n=1,expand = True)
2
3 df['CustomerId'] = new_df[1].astype(int)
```

In [28]:

```
1 new_df = df['TransactionId'].str.split('_',n=1,expand = True)
2
3 df['TransactionId'] = new_df[1].astype(int)
```

In [29]:

```
1 new_df = df['AccountId'].str.split('_',n=1,expand = True)
2
3 df['AccountId'] = new_df[1].astype(int)
```

In [30]:

```
1 new_df = df['SubscriptionId'].str.split('_',n=1,expand = True)
2
3 df['SubscriptionId'] = new_df[1].astype(int)
```

In [31]:

```
1 new_df = df['ProviderId'].str.split('_',n=1,expand = True)
2
3 df['ProviderId'] = new_df[1].astype(int)
```

In [32]:

```
1 new_df = df['ProductId'].str.split('_',n=1,expand = True)
2
3 df['ProductId'] = new_df[1].astype(int)
```

In [33]:

```
1 new_df = df['ChannelId'].str.split('_',n=1,expand = True)
2
3 df['ChannelId'] = new_df[1].astype(int)
```


In [45]:

```
1 df
```

Out[45]:

	TransactionId	BatchId	AccountId	SubscriptionId	CustomerId	CurrencyCode	CountryCode
0	76871	36123	3957	887	36123	UGX	
1	73770	15642	4841	3829	15642	UGX	
2	26203	53941	4229	222	53941	UGX	
3	380	102363	648	2185	102363	UGX	
4	28195	38780	4841	3829	38780	UGX	
...
95657	89881	96668	4841	3829	96668	UGX	
95658	91597	3503	3439	2643	3503	UGX	
95659	82501	118602	4841	3829	118602	UGX	
95660	136354	70924	1346	652	70924	UGX	
95661	35670	29317	4841	3829	29317	UGX	

95662 rows × 8 columns



In [47]:

```
1 df.drop(['CurrencyCode','CountryCode'], axis = 1, inplace = True)
2 df.head()
```

Out[47]:

	TransactionId	BatchId	AccountId	SubscriptionId	CustomerId	ProviderId	ProductId	Product
0	76871	36123	3957	887	36123	6	10	
1	73770	15642	4841	3829	15642	4	6	financi
2	26203	53941	4229	222	53941	6	1	
3	380	102363	648	2185	102363	1	21	
4	28195	38780	4841	3829	38780	4	6	financi



In [48]:

```

1 #getting dummies for some columns that are categorical
2 selected_categorical_columns = ["ProviderId", "ProductId", "ProductCategory", "ChannelId"]
3 new_df = pd.get_dummies(df, columns= selected_categorical_columns, drop_first = True)
4 new_df.head()

```

Out[48]:

	TransactionId	BatchId	AccountId	SubscriptionId	CustomerId	Amount	Value	FraudResult
0	76871	36123	3957	887	36123	1000.0	1000	0
1	73770	15642	4841	3829	15642	-20.0	20	0
2	26203	53941	4229	222	53941	500.0	500	0
3	380	102363	648	2185	102363	20000.0	21800	0
4	28195	38780	4841	3829	38780	-644.0	644	0

In [49]:

```
1 new_df.shape
```

Out[49]:

(95662, 52)

In [50]:

```

1 #split data into target and feature
2
3 X = new_df.drop(['FraudResult'], axis =1)
4 y= new_df['FraudResult']

```

In [72]:

```

1 #split into test and train data
2
3 from sklearn.model_selection import train_test_split
4
5 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=

```

In [73]:

```

1
2 X_train.shape, X_test.shape

```

Out[73]:

((76529, 51), (19133, 51))

In [74]:

```
1 X_train
```

Out[74]:

	TransactionId	BatchId	AccountId	SubscriptionId	CustomerId	Amount	Value	Year	M
66339	140085	114864	4841	3829	114864	-50.0	50	2019	
87279	64558	39429	3981	910	39429	1000.0	1000	2019	
40582	134904	1673	135	3595	1673	15000.0	16650	2018	
58655	95030	133112	4840	3829	133112	-1000.0	1000	2019	
87335	75383	7649	4841	3829	7649	-50.0	50	2019	
...
21243	1505	124507	575	4369	124507	2000.0	2000	2018	
45891	68325	94167	4841	3829	94167	-50.0	50	2018	
42613	62369	85767	571	873	85767	1500.0	1500	2018	
43567	1639	36597	2123	1456	36597	2000.0	2000	2018	
68268	83620	99997	2659	3327	99997	1000.0	1000	2019	

76529 rows × 51 columns



In [71]:

```
1 X_test
```

Out[71]:

```
array([[0.64961075, 0.84289658, 0.17210744, ..., 0.      , 0.      ,
        1.      ],
       [0.84885002, 0.17318516, 0.85247934, ..., 0.      , 1.      ,
        0.      ],
       [0.88152572, 0.36934935, 0.03904959, ..., 0.      , 1.      ,
        0.      ],
       ...,
       [0.96174327, 0.23050726, 0.6142562 , ..., 0.      , 1.      ,
        0.      ],
       [0.78433045, 0.4117297 , 0.54690083, ..., 0.      , 1.      ,
        0.      ],
       [0.75759838, 0.49365119, 0.50454545, ..., 0.      , 1.      ,
        0.      ]])
```


In [56]:

```
1 X_train.dtypes
```

Out[56]:

```
TransactionId      int32
BatchId            int32
AccountId          int32
SubscriptionId     int32
CustomerId         int32
Amount            float64
Value             int64
Year              int64
Month             int64
day              int64
ProviderId_2       uint8
ProviderId_3       uint8
ProviderId_4       uint8
ProviderId_5       uint8
ProviderId_6       uint8
ProductId_2        uint8
ProductId_3        uint8
ProductId_4        uint8
```

In [76]:

```
1 #feature scaling
2
3 from sklearn.preprocessing import MinMaxScaler
4 scaler= MinMaxScaler()
5
6 X_train = scaler.fit_transform(X_train)
7 X_test = scaler.fit_transform(X_test)
```

In [63]:

```
1 X_train
```

Out[63]:

```
array([[0.99576343, 0.82343664, 1.          , ..., 0.          , 1.          ,
        0.          ],
       [0.45889252, 0.28264906, 0.82231405, ..., 0.          , 0.          ,
        1.          ],
       [0.95893517, 0.01197927, 0.02768595, ..., 0.          , 1.          ,
        0.          ],
       ...,
       [0.44333239, 0.61484253, 0.1177686 , ..., 0.          , 1.          ,
        0.          ],
       [0.01164345, 0.26234667, 0.43842975, ..., 0.          , 1.          ,
        0.          ],
       [0.59439153, 0.71685628, 0.54917355, ..., 0.          , 1.          ,
        0.          ]])
```

In [64]:

```
1 X_test
```

Out[64]:

```
array([[0.64961075, 0.84289658, 0.17210744, ..., 0.        , 0.        ,
        1.        ],
       [0.84885002, 0.17318516, 0.85247934, ..., 0.        , 1.        ,
        0.        ],
       [0.88152572, 0.36934935, 0.03904959, ..., 0.        , 1.        ,
        0.        ],
       ...,
       [0.96174327, 0.23050726, 0.6142562 , ..., 0.        , 1.        ,
        0.        ],
       [0.78433045, 0.4117297 , 0.54690083, ..., 0.        , 1.        ,
        0.        ],
       [0.75759838, 0.49365119, 0.50454545, ..., 0.        , 1.        ,
        0.        ]])
```

In [75]:

```
1 cols = X_train.columns
```


In [79]:

```
1 X_test = pd.DataFrame(X_test, columns=[cols])
2 X_test
```

Out[79]:

	TransactionId	BatchId	AccountId	SubscriptionId	CustomerId	Amount	Value	Yea
0	0.649611	0.842897	0.172107	0.777203	0.842897	0.092258	0.000203	1.
1	0.848850	0.173185	0.852479	0.876707	0.173185	0.092176	0.000112	1.
2	0.881526	0.369349	0.039050	0.392015	0.369349	0.092258	0.000203	0.
3	0.086950	0.909783	0.999793	0.791684	0.909783	0.091981	0.000101	1.
4	0.192229	0.127887	0.326033	0.169425	0.127887	0.092166	0.000101	1.
...
19128	0.948576	0.287148	0.820248	0.320025	0.287148	0.092258	0.000203	1.
19129	0.548726	0.905682	1.000000	0.791684	0.905682	0.091613	0.000507	1.
19130	0.961743	0.230507	0.614256	0.476210	0.230507	0.092248	0.000192	1.
19131	0.784330	0.411730	0.546901	0.270790	0.411730	0.092166	0.000101	1.
19132	0.757598	0.493651	0.504545	0.729003	0.493651	0.092258	0.000203	0.

19133 rows × 51 columns



In []:

```
1
```

In [78]:

```
1 X_train = pd.DataFrame(X_train, columns=[cols])
2 X_train
```

Out[78]:

	TransactionId	BatchId	AccountId	SubscriptionId	CustomerId	Amount	Value	Yea
0	0.995763	0.823437	1.000000	0.791684	0.823437	0.083483	0.000005	1.
1	0.458893	0.282649	0.822314	0.187836	0.282649	0.083581	0.000101	1.
2	0.958935	0.011979	0.027686	0.743277	0.011979	0.084879	0.001685	0.
3	0.675498	0.954255	0.999793	0.791684	0.954255	0.083395	0.000101	1.
4	0.535840	0.054821	1.000000	0.791684	0.054821	0.083483	0.000005	1.
...
76524	0.010691	0.892567	0.118595	0.903393	0.892567	0.083673	0.000202	0.
76525	0.485670	0.675061	1.000000	0.791684	0.675061	0.083483	0.000005	0.
76526	0.443332	0.614843	0.117769	0.180182	0.614843	0.083627	0.000152	0.
76527	0.011643	0.262347	0.438430	0.300786	0.262347	0.083673	0.000202	0.
76528	0.594392	0.716856	0.549174	0.687836	0.716856	0.083581	0.000101	1.

76529 rows × 51 columns

In [80]:

```
1 #Model training
2 from sklearn.linear_model import LogisticRegression
3
4 logreg = LogisticRegression(solver='liblinear', random_state=0)
5
6 logreg.fit(X_train, y_train)
```

Out[80]:

LogisticRegression(random_state=0, solver='liblinear')

In [81]:

```
1 #test data
2 y_pred_test = logreg.predict(X_test)
```

In [82]:

```
1 y_pred_test
```

Out[82]:

array([0, 0, 0, ..., 0, 0, 0], dtype=int64)

In [83]:

```
1 #probability of getting zero, i.e , no fraud
2
3 logreg.predict_proba(X_test)[: ,0]
```

Out[83]:

```
array([0.99753105, 0.99977752, 0.99987278, ..., 0.99418927, 0.99983678,
       0.9998944 ])
```

In [84]:

```
1 #check for accuracy
2
3 from sklearn.metrics import accuracy_score
4
5 print('Model accuracy score(test): ', accuracy_score(y_test, y_pred_test))
```

Model accuracy score(test): 0.998484294151466

In [85]:

```
1 y_pred_train = logreg.predict(X_train)
2
3 print('Model accuracy score(train): ', accuracy_score(y_train, y_pred_train))
```

Model accuracy score(train): 0.99820982895373

In [86]:

```
1 #model accuracy metrics for logistic regression
2 print('Training set score: ', logreg.score(X_train, y_train))
3
4 print('Test set score: ', logreg.score(X_test, y_test))
```

Training set score: 0.99820982895373

Test set score: 0.998484294151466

In [87]:

```
1 #check null accuracy
2 null_acuracy = 22726/(len(y_test))
```

In [88]:

```
1 null_acuracy
```

Out[88]:

1.1877907280614646

In [89]:

```
1 from sklearn.metrics import confusion_matrix
2 cm =confusion_matrix(y_test, y_pred_test)
```

In [90]:

```
1 ca
```

Out[90]:

```
array([[19098,    0],
       [   29,    6]], dtype=int64)
```

In [91]:

```
1 print('True positives (TP) = ', cm[0,0])
2 print('False positives (FP)= ', cm[0, 1])
3 print('False Negatives (FN)= ', cm[1, 0])
4 print('True Negatives (TN))= ', cm[1, 1])
```

```
True positives (TP) = 19098
False positives (FP)= 0
False Negatives (FN)= 29
True Negatives (TN))= 6
```

In [92]:


```
1 #visualze with heatmap
2 from sklearn.metrics import ConfusionMatrixDisplay
3 disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=logreg.classes_)
```

In [93]:

```
1 disp.plot()
```

Out[93]:

```
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x17aff87c490>
```



In [94]:

```
1 from sklearn.metrics import classification_report
2
3 #classification_report(y_test, y_pred_test)
```

In [95]:

```
1 #new_df
```

In [42]:

```
1 #new_df2 = new_df.drop(['CurrencyCode', 'CountryCode'], axis = 1, inplace=True)
```

In [44]:

```
1 #print(new_df2)
```

None

In []:

```
1 #split data into target and feature
2
3 #X = new_df.drop(['CurrencyCode', 'CountryCode', 'FraudResult'], axis =1)
4 #y= new_df['FraudResult']
```

In [96]:

```
1 print(classification_report(y_test, y_pred_test))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	19098
1	1.00	0.17	0.29	35
accuracy			1.00	19133
macro avg	1.00	0.59	0.65	19133
weighted avg	1.00	1.00	1.00	19133

In []:

```
1
```

In []:

```
1
```

In []:

```
1
```

In []:

```
1
```

In []:

```
1
```

In []:

```
1
```

In []:

```
1 #split into test and train data
2
3 from sklearn.model_selection import train_test_split
4
5 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=
```

In []:

```
1 #check the shapeof our split
2
3 #X_train.shape, X_test.shape
4
```

In []:

```
1
```

In []:

```
1
```

In []:

```
1
2 #X_train = pd.concat([X_train,ProductCategory_dummy, productId_dummy,providerId_dummy,C
3
```

In []:

```
1
```

In []:

```
1
```

In []:

```
1
```

In []:

```
1
```

In []:

```
1
```

In []:

```
1
```

In []:

```
1
```

In []:

1	
---	--

In []:

1	
---	--

In []:

1	
---	--

In []:

1	
---	--

In []:

1	
---	--

In []:

1	
---	--

In []:

1	
---	--

In []:

1	
---	--

In []:

1	
---	--

In []:

1	
---	--

In []:

1	
---	--

In []:

1	
---	--

In []:

1	
---	--

In []:

1	
---	--

In []:

1

In []:

1

In []:

1

In []:

1

In []:

1

In []:

1

In []:

1

In []:

1

In []:

1

In []:

1